

Testing Linear Hypotheses

Lecture 4 (Stata code)

Table of contents

Reading	2
Hypothesis Tests	2
Simple hypotheses	2
Linear combinations	3
Multiple linear combinations	3
Single linear combinations	3
Example: PM2.5	4
Null hypotheses	4
Alternative hypothesis	5
Test statistic	5
Variance	6
Distribution	7
Example: PM2.5	7
Multiple linear hypotheses	7
Example: PM2.5	8
Test	8
Test statistic	9
Test statistic	9
Distribution	10
Rejection rule	10
Alternative form	11
Example: PM2.5	11
Revisiting single linear restriction	11
Example: PM2.5	12
Regression significance	12
Stata output	13

Structural Change	14
Dummy variables: recap	14
Dummy variable interactions	14
Example	15
F -test	16
Test:	16
Null hypothesis	17
Approach 1	17
Distribution	17
Approach 2	18
Interactions	18
Test	19
F -stat	19
Equivalence	19
Demonstration	20
References	21
Bibliography	21

Reading

This lecture covers material from Wooldridge (2025):

- Chapter 4.4 - “Testing Hypotheses about a Single Linear Combination of the Parameters”
- Chapter 4.5 - “Testing Multiple Linear Restrictions”
- Chapter 7.4 - “Interactions Involving Dummy Variables”

Note: I have edited the syllabus slightly to move the topic of heteroskedasticity to Week 5 with asymptotics.

Hypothesis Tests

Simple hypotheses

So far, we have considered hypotheses that involve a single estimator:

1. $H_0 : \beta_j = \kappa_0$
2. $H_0 : \beta_j \geq \kappa_0$
3. $H_0 : \beta_j \leq \kappa_0$

Linear combinations

Today, we consider hypotheses of the form:

$$H_0 : a_0\beta_0 + a_1\beta_1 + \cdots + a_k\beta_j = \kappa_0$$

for constants $\{a_j\}_{j=0,1,\dots,k}$. This is a **linear combination of parameters**.

For example, $H_0 : \beta_1 = 2\beta_2 - 0.5$ can be written as,

$$H_0 : \underbrace{0}_{a_0} \beta_0 + \underbrace{1}_{a_1} \beta_1 + \underbrace{(-2)}_{a_2} \beta_2 = \underbrace{-0.5}_{\kappa_0}$$

Multiple linear combinations

Similarly, we could have q of these statements,

$$\begin{aligned} H_0 : a_0^1\beta_0 + a_1^1\beta_1 + \cdots + a_k^1\beta_j &= \kappa_0^1 \\ a_0^2\beta_0 + a_1^2\beta_1 + \cdots + a_k^2\beta_j &= \kappa_0^2 \\ &\vdots \\ a_0^q\beta_0 + a_1^q\beta_1 + \cdots + a_k^q\beta_j &= \kappa_0^q \end{aligned}$$

Single linear combinations

Consider the following example related to environmental and health economics

$$\ln(PM2.5_\ell) = \beta_0 + \beta_1 EV_\ell + \beta_2 hybrid_\ell + \mathbf{x}_i\gamma + u_i$$

where,

- $PM2.5$ is micrograms of particulate matter (with diameter ≤ 2.5 micrometres) per cubic metre in location ℓ .
 - standard measure for health effects as absorbable by lungs and bloodstream
- EV and $hybrid$ are number of registered vehicles (/1000) by fuel type in location ℓ .

Example: PM2.5

The following estimates use local authority level data for 2019 in the UK.

```
* Open CSV data
import delimited "$data_dir\pm25.csv", clear

* Estimate univariate linear regression
reg ln_pm25 ev_k hybrid_k petrol_k diesel_k ln_pop
```

(encoding automatically selected: ISO-8859-1)
(17 vars, 291 obs)

Source	SS	df	MS	Number of obs	=	291
-----+-----				F(5, 285)	=	33.45
Model	2.84542086	5	.569084172	Prob > F	=	0.0000
Residual	4.84835963	285	.017011788	R-squared	=	0.3698
-----+-----				Adj R-squared	=	0.3588
Total	7.69378049	290	.026530278	Root MSE	=	.13043

ln_pm25	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
-----+-----						
ev_k	.0779783	.03734	2.09	0.038	.0044811	.1514755
hybrid_k	.0322276	.0063524	5.07	0.000	.019724	.0447312
petrol_k	.0000559	.0007053	0.08	0.937	-.0013323	.0014442
diesel_k	-.003546	.0005762	-6.15	0.000	-.0046802	-.0024119
ln_pop	.088135	.0208428	4.23	0.000	.0471096	.1291604
_cons	1.264753	.2325884	5.44	0.000	.8069436	1.722562

Null hypotheses

From a public policy perspective (if you consider designing subsidies), you want to estimate both β_1 and β_2 , BUT also know if

$$H_0 : \beta_1 = \beta_2$$

That is, should the policy discriminate between electrified fuel types?

- Engineers may be able provide mechanical differences, but they can't tell you about differences in usage (driver behaviour) which is important for realized pollution mitigation.

Warning

Do not take these estimates too seriously. This model is a classic example of selection into “treatment”: those who live in higher pollution areas (e.g. dense urban centres) are also more likely to buy EVs. A better research design would use panel data to evaluate this question dynamically.

Alternative hypothesis

The alternative in this case, may be

$$H_1 : \beta_1 \neq \beta_2$$

But may also be,

$$H_1 : \beta_1 < \beta_2$$

or

$$H_1 : \beta_1 > \beta_2$$

if there is a strong case for ruling out one case (see Wooldridge 2025, 139).

Test statistic

For $H_0 : \beta_1 = \beta_2 \Rightarrow H_0 : \beta_1 - \beta_2 = 0$

$$\text{test-stat} = \frac{\hat{\beta}_1 - \hat{\beta}_2}{\sqrt{\text{Var}(\hat{\beta}_1 - \hat{\beta}_2)}}$$

Note

The more general case would be $H_0 : a_1\beta_1 + a_2\beta_2 = \kappa_0$

$$\text{test-stat} = \frac{a_1\hat{\beta}_1 + a_2\hat{\beta}_2 - \kappa_0}{\sqrt{\text{Var}(a_1\hat{\beta}_1 + a_2\hat{\beta}_2)}}$$

which can be extended to include more parameters.

Variance

The key to understanding and executing this test is,¹

$$\text{Var}(\hat{\beta}_1 - \hat{\beta}_2) \neq \text{Var}(\hat{\beta}_1) + \text{Var}(\hat{\beta}_2)$$

This is because,

$$\begin{aligned}\text{Var}(a_1z_1 + a_2z_2) &= \text{Var}(a_1z_1) + \text{Var}(a_2z_2) + 2 \times \text{Cov}(a_1z_1, a_2z_2) \\ &= a_1^2\text{Var}(z_1) + a_2^2\text{Var}(z_2) + 2 \times a_1a_2\text{Cov}(z_1, z_2)\end{aligned}$$

for random variables z_1 and z_2 and constants a_1 and a_2 .

So,

$$\text{Var}(\hat{\beta}_1 - \hat{\beta}_2) = \text{Var}(\hat{\beta}_1) + \text{Var}(\hat{\beta}_2) - \underbrace{2 \times \text{Cov}(\hat{\beta}_1, \hat{\beta}_2)}_{\neq 0}$$

Why -2 ?

- For this H_0 ,

$$\Rightarrow H_0 : \underbrace{\beta_1 - \beta_2 = 0}_{a_1\beta_1 + a_2\beta_2 = \kappa_0}$$

So, $a_1 = 1$ and $a_2 = -1$.

¹We drop the $\text{Var}(\cdot|X)$ notation from [Lecture 3](#). However, all these variances are conditional on the sample values of the independent variables.

Distribution

Under MLR 1-6, the **conditional** distribution of this test-statistic is t_{n-k-1}

$$t\text{-stat} = \frac{\hat{\beta}_1 - \hat{\beta}_2}{\sqrt{\text{Var}(\hat{\beta}_1 - \hat{\beta}_2)}} \underset{H_0}{\sim} t_{n-k-1}$$

under the null.

- You can use the same steps from [Lecture 3](#) to conduct inference: p -value, critical value.
- We tend not to construct confidence intervals in such cases.
- Similarly, you can compute the power under the alternative $\beta_1 - \beta_2 = \kappa$.

Example: PM2.5

```
* t-test  
lincom ev_k - hybrid_k
```

```
( 1)  ev_k - hybrid_k = 0
```

ln_pm25		Coefficient	Std. err.	t	P> t	[95% conf. interval]
-----+-----						
(1)		.0457507	.0423487	1.08	0.281	-.0376053 .1291066

Multiple linear hypotheses

Another common form of hypothesis concerns **multiple** linear combinations of estimators.

$$\begin{aligned} H_0 : a_0^1\beta_0 + a_1^1\beta_1 + \dots + a_k^1\beta_j &= \kappa_0^1 \\ a_0^2\beta_0 + a_1^2\beta_1 + \dots + a_k^2\beta_j &= \kappa_0^2 \\ &\vdots \\ a_0^q\beta_0 + a_1^q\beta_1 + \dots + a_k^q\beta_j &= \kappa_0^q \end{aligned}$$

💡 Linear transformation

You can think of this as a linear transformation of β .

$$H_0 : R\beta = r$$

Where R is a $q \times k$ matrix of constants and r a q -dimensional column vector of constants.

Example: PM2.5

In our example of

$$PM2.5_\ell = \beta_0 + \beta_1 EV_\ell + \beta_2 hybrid_\ell + \mathbf{x}_i \gamma + u_i$$

This could be as simple as,

$$H_0 : \beta_1 = 0 \text{ and } \beta_2 = 0$$

🔥 Multiple hypothesis testing

Why not test $H_0^{(1)} : \beta_1 = 0$ and then $H_0^{(2)} : \beta_2 = 0$ as separate individual hypotheses? The combined size of these tests would not be α . Why? The estimators are correlated which implies:

$$Pr(\text{Reject } H_0^{(1)} \text{ or } H_0^{(2)} | \beta_1 = 0, \beta_2 = 0) \neq \alpha$$

💡 Multiple hypothesis testing

There are a number of ways to correctly do multiple hypothesis testing. One is the [Bonferroni correction](#) which gets you to reject separate hypotheses sequentially using an algorithm. It does so by ensuring that the family-wise error rate (the probability of rejecting at least one true H_0 across a “family” of hypotheses) is $\leq \alpha$.

Test

The test is performed by considering two models:

1. **Unrestricted** model; using our example,

$$PM2.5_\ell = \beta_0 + \beta_1 EV_\ell + \beta_2 hybrid_\ell + \mathbf{x}_i \gamma + u_i$$

- storing the SSR_{ur} from this model
2. **Restricted model**, which imposes $H_0 : \beta_1 = 0$ and $\beta_2 = 0$,

$$PM2.5_\ell = \beta_0 + \mathbf{x}_i \gamma + u_i$$

- storing the SSR_r from this model

Test statistic

The test static that is then given by,

$$F\text{-stat} = \frac{(SSR_r - SSR_{ur})/q}{SSR_{ur}/df_{ur}}$$

where,

- SSR_r and SSR_{ur} come from the restricted (r) and unrestricted models (ur).²

$$SSR_r - SSR_{ur} > 0$$

Test statistic

The test static that is then given by,

$$F\text{-stat} = \frac{(SSR_r - SSR_{ur})/q}{SSR_{ur}/df_{ur}}$$

and,

- q is the number of restrictions (i.e. equations tested) and is given by the difference in residual degrees of freedom

$$q = df_r - df_{ur}$$

²This inequality is how you can remember the order. The F -stat is strictly positive.

Distribution

Under the null, F -stat has an F distribution (hence the name)

$$F\text{-stat} \underset{H_0}{\sim} F_{q,n-k-1}$$

We reject H_0 if

$$F\text{-stat} > F_{q,n-k-1}^{-1}(1 - \alpha) = f_{q,n-k-1,1-\alpha}$$

where $F_{q,n-k-1}^{-1}(1 - \alpha)$ is the $1 - \alpha$ percentile of the distribution.

Rejection rule

$$F\text{-stat} > f_{q,n-k-1,1-\alpha}$$

- **Always a one-sided test.**³
- Doesn't mean that H_1 is one-sided.
 - You cannot tell which direction the deviation from the null is, because F -stat uses squared-deviations (i.e. sign is lost).
- $H_1 : H_0$ false, or $H_1 : \beta_1 \neq 0$ or $\beta_2 \neq 0$

F distribution

The F distribution is defined as the ratio of two chi-squared distributions,

$$F_{q,m} = \frac{\chi_q^2/q}{\chi_m^2/m}$$

each normalized by their corresponding degrees of freedom.

³This does not mean that F -distributed tests are always one-sided, but rather that multiple linear hypotheses of this form are.

Alternative form

Since $SSR = (1 - \mathbf{R}^2)SST$ and $SST_r = SST_{ur}$,

We can write the F -stat as,

$$F\text{-stat} = \frac{(\mathbf{R}_{ur}^2 - \mathbf{R}_r^2)/q}{(1 - \mathbf{R}_{ur}^2)/(n - k - 1)}$$

Note,

- The order switches in the numerator, but remains positive since $\mathbf{R}_{ur}^2 > \mathbf{R}_r^2$
- $n - k - 1 = df_{ur}$, with k the number of parameters in the unrestricted model

Example: PM2.5

```
* F-test  
test ev_k hybrid_k
```

```
( 1)  ev_k = 0  
( 2)  hybrid_k = 0  
  
F( 2, 285) = 53.88  
Prob > F = 0.0000
```

Revisiting single linear restriction

You can think of a single linear hypothesis as a special case of multiple where $q = 1$. For example,

$$H_0 : \beta_1 = \beta_2$$

- Restricted model:

$$\begin{aligned} PM2.5_\ell &= \beta_0 + \beta_1 EV_\ell + \underbrace{\beta_1}_{=\beta_2} hybrid_\ell + \mathbf{x}_i \gamma + u_i \\ &= \beta_0 + \beta_1 (EV_\ell + hybrid_\ell) + \mathbf{x}_i \gamma + u_i \end{aligned}$$

A regression with a new transformed variable: $EV_\ell + hybrid_\ell$

Moreover, you can show that for $q = 1$

$$F\text{-stat} = \frac{SSR_r - SSR_{ur}}{SSR_{ur}/(n - k - 1)} = (t\text{-stat})^2$$

Example: PM2.5

```
* t-test
lincom ev_k - hybrid_k

* F-test
test ev_k=hybrid_k

* check
dis sqrt(r(F))
```

```
( 1)  ev_k - hybrid_k = 0
```

ln_pm25	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
(1)	.0457507	.0423487	1.08	0.281	-.0376053	.1291066

```
( 1)  ev_k - hybrid_k = 0
```

```
F( 1, 285) = 1.17
Prob > F = 0.2809
```

```
1.0803322
```

Regression significance

A special case of this form of test concerns the joint significance of all the *slope* parameters in a model:

$$H_0 : \beta_1 = 0, \beta_2 = 0, \dots, \beta_k = 0$$

In this case,

- The restricted model is a regression on a constant (alone) where $\Rightarrow \mathbf{R}_r^2 = 0$
- And $\mathbf{R}_{ur}^2$ is simply the $\mathbf{R}^2$ from the model
- The number of restrictions is $q = k$

We define the F -stat as,

$$F\text{-stat} = \frac{\mathbf{R}^2/k}{(1 - \mathbf{R}^2)/(n - k - 1)}$$

Stata output

This test is reported as standard with the Stata output as:

- $F(,)$: F -stat
- Prob > F: corresponding p -value of F -test

Source	SS	df	MS	Number of obs	=	291
				F(5, 285)	=	33.45
Model	2.84542086	5	.569084172	Prob > F	=	0.0000
Residual	4.84835963	285	.017011788	R-squared	=	0.3698
				Adj R-squared	=	0.3588
Total	7.69378049	290	.026530278	Root MSE	=	.13043

ln_pm25	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
ev_k	.0779783	.03734	2.09	0.038	.0044811	.1514755
hybrid_k	.0322276	.0063524	5.07	0.000	.019724	.0447312
petrol_k	.0000559	.0007053	0.08	0.937	-.0013323	.0014442
diesel_k	-.003546	.0005762	-6.15	0.000	-.0046802	-.0024119
ln_pop	.088135	.0208428	4.23	0.000	.0471096	.1291604
_cons	1.264753	.2325884	5.44	0.000	.8069436	1.722562

Structural Change

Dummy variables: recap

Recall, when we have the model:

$$y_i = \beta_0 + \beta_1 d_i + u_i$$

where $d_i \in \{0, 1\}$. Then,

$$\beta_1 = E[y_i | d_i = 1] - E[y_i | d_i = 0]$$

Dummy variable interactions

Now, suppose the model was

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 d_i + \beta_3 x_i \times d_i + u_i$$

where x_i is continuous.

- What is the interpretation of each parameter?

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 d_i + \beta_3 x_i \times d_i + u_i$$

- $\beta_0 = E[y_i | d_i = 0, x_i = 0]$
- $\beta_1 = \frac{dE[y_i | d_i = 0, x_i]}{dx_i}$
- $\beta_2 = E[y_i | d_i = 1, x_i = 0] - E[y_i | d_i = 0, x_i = 0]$
- $\beta_3 = \frac{dE[y_i | d_i = 1, x_i]}{dx_i} - \frac{dE[y_i | d_i = 0, x_i]}{dx_i}$

Another way of thinking about it:

- For $d_i = 0$, the model is:

$$y_i = \beta_0 + \beta_1 x_i + u_i$$

- For $d_i = 1$, the model is:

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 + \beta_3 x_i + u_i$$

Which can be rewritten as:

$$y_i = \underbrace{(\beta_0 + \beta_2)}_{\gamma_0} + \underbrace{(\beta_1 + \beta_3)}_{\gamma_1} x_i + u_i$$

We have two models: for $d_i = 0$ and $d_i = 1$.

- β_2 is the difference in intercept
- β_3 is the difference in slope

Example

The typical example is the returns to education for men and women (as in Wooldridge 2025, 243)

- But, given what we know about the ‘parent penalty’ (Kleven et al. 2019), it seems more pertinent to focus on experience

$$\ln(wage_i) = \beta_0 + \beta_1 exp_i + u_i$$

```
use "$data_dir\morg24_bothgen.dta", clear

* Scatter
twoway (scatter lnwage exp, mc(blue%50)) (lfit lnwage exp, lwidth(*2)), by(sex, legend(off))
```

If we overlay the fitted lines we get

```
* overlay fitted lines
twoway (lfit lnwage exp if sex==1, lc(blue)) (lfit lnwage exp if sex==2, lc(red) lwidth(*2))
```

***F*-test**

The *F*-test can be used to test differences in a model across groups.⁴

- For example, whether the Mincer equation (from [Tutorial 1](#)) differs by gender.

$$\ln(wage_i) = \beta_0 + \beta_1 edu_i + \beta_2 exp_i + u_i$$

- We estimated this model for the male-only sample
- The quadratic term in *exp* has been dropped for simplification
- The model has $k + 1 = 3$ parameters

Test:

We may wish to ask whether:

- the model for men (group = 1; e.g. “Male”)

$$\ln(wage_i) = \beta_0 + \beta_1 edu_i + \beta_2 exp_i + u_i$$

is the same as,

- the model for women (group = 2; e.g. “Female”)?

$$\ln(wage_i) = \gamma_0 + \gamma_1 edu_i + \gamma_2 exp_i + u_i$$

⁴We will look at the case of two groups, but the approach can be extended to multiple groups.

Null hypothesis

This question implies a null hypothesis:

$$H_0 : \beta_0 = \gamma_0, \beta_1 = \gamma_1, \beta_2 = \gamma_2$$

against,

$$H_1 : H_0 \text{ false}$$

Approach 1

The F -test for this hypothesis has the form:

$$F\text{-stat} = \frac{(SSR_P - (SSR_1 + SSR_2))/(k+1)}{(SSR_1 + SSR_2)/(n - 2(k+1))}$$

where,

- SSR_P is the “pooled” (P) SSR from a regression of the joint sample.
- SSR_1 is the SSR from a regression for just group 1.
- SSR_2 is the SSR from a regression for just group 2.
 - together SSR_1 and SSR_2 represent the unrestricted model.
- $n - 2(k+1)$ is the **combined** degrees of freedom from group 1 and 2.
 - $k+1$ parameters are estimated in each group.

Distribution

The F -stat is (unsurprisingly) F -distributed,

$$F\text{-stat} \underset{H_0}{\sim} F_{k+1, n-2(k+1)}$$

Approach 2

This test has an equivalent form using interactions. Let,

$$d_i = \mathbf{1}\{\text{'Group 2'}\}$$

For example,

$$d_i = \mathbf{1}\{\text{'Female'}\}$$

Interactions

We can write the unrestricted model as,

$$\ln(wage_i) = \beta_0 + \beta_1 edu_i + \beta_2 exp_i + \beta_3 d_i + \beta_4 edu_i \times d_i + \beta_5 exp_i \times d_i + u_i$$

We now have 6 parameters: $2 \times (k + 1)$ (for $k = 2$)

As before, the dummy and dummy-interactions tell you about differences,

$$\beta_3 = E[\ln(wage_i)|\textcolor{red}{d_i} = \textcolor{red}{1}, edu_i = 0, exp_i = 0] - E[\ln(wage_i)|\textcolor{blue}{d_i} = \textcolor{blue}{0}, edu_i = 0, exp_i = 0]$$

and

$$\beta_j = \frac{\partial E[\ln(wage_i)|\textcolor{red}{d_i} = \textcolor{red}{1}, edu_i, exp_i]}{\partial x_{ij}} - \frac{\partial E[\ln(wage_i)|\textcolor{blue}{d_i} = \textcolor{blue}{0}, edu_i, exp_i]}{\partial x_{ij}} \quad j = 4, 5$$

Test

Intuitively, the test for whether the model differs by group is given by,

$$H_0 : \beta_3 = 0, \beta_4 = 0, \beta_5 = 0$$

where,

- **unrestricted** model is given by,

$$\ln(wage_i) = \beta_0 + \beta_1 edu_i + \beta_2 exp_i + \beta_3 d_i + \beta_4 edu_i \times d_i + \beta_5 exp_i \times d_i + u_i$$

- $df_{ur} = n - 6$

- **restricted** model is given by,

$$\ln(wage_i) = \beta_0 + \beta_1 edu_i + \beta_2 exp_i + v_i$$

- $df_{ur} = n - 3$

F-stat

The F -stat is given by,

$$F\text{-stat} = \frac{(SSR_r - SSR_{ur})/3}{SSR_{ur}/(n - 6)}$$

Compare this to the first approach:

$$F\text{-stat} = \frac{(SSR_P - (SSR_1 + SSR_2))/3}{(SSR_1 + SSR_2)/(n - 6)}$$

Equivalence

These two test-statistics are equivalent because,

1. $SSR_P = SSR_r$
2. $SSR_1 + SSR_2 = SSR_{ur}$

Demonstration

Here are the SSR for each version:

- Model for group 1 (i.e. “male”)

```
qui reg lnwage edu exp if sex==1  
dis e(rss)
```

715.90864

- Model for group 2 (i.e. “female”)

```
qui reg lnwage edu exp if sex==2  
dis e(rss)
```

539.59768

- Unrestricted model (i.e. with interactions)

```
qui reg lnwage i.sex##c.edu i.sex##c.exp  
dis e(rss)
```

1255.5063

- Restricted model (i.e. without interactions)

```
qui reg lnwage edu exp  
dis e(rss)
```

1307.1667

Here is how you would programme the test

- Estimate interacted model

```
qui reg lnwage i.sex#c.edu i.sex#c.exp  
test 2.sex 2.sex#c.edu 2.sex#c.exp
```

```
( 1)  2.sex = 0  
( 2)  2.sex#c.edu = 0  
( 3)  2.sex#c.exp = 0
```

```
      F( 3, 4994) =    68.50  
      Prob > F =    0.0000
```

The notation `2.sex` is how Stata names the coefficient on level 2 of the factor variable `sex`, and `2.sex#c.edu` is the interaction with `c.edu` (a continuous variable).

References

Bibliography

Kleven, Henrik, Camille Landais, Johanna Posch, Andreas Steinhauer, and Josef Zweimüller. 2019. “Child Penalties Across Countries: Evidence and Explanations.” In *Aea Papers and Proceedings*, 109:122–26. American Economic Association 2014 Broadway, Suite 305, Nashville, TN 37203.

Wooldridge, Jeffrey M. 2025. “Introductory Econometrics: A Modern Approach.”